

Contents

- 1 Selecting the dataset
 - 1.1 Plot chosen dataset
- 2 Stationary Test
 - 2.1 ADF Test with trend
- 3 Mitigation of non-stationarity
 - 3.1 ADF Test
 - 3.2 KPSS test
- 4 Parameter selection for time series modeling
 - 4.1 ACF and PACF plots of the chosen datasets (First difference)
- 5 Modeling the series
 - 5.1 ARIMA(1,1,2)
 - 5.2 Auto-ARIMA
 - 5.2.1 True value vs fitted value
 - 5.2.2 Model diagnostics
 - 5.3 Modeling heteroskedasticity in residuals with GARCH
 - 5.3.1 ACF and PACF of the residuals and the squares
 - 5.3.2 GARCH fit
 - 5.3.3 Diagnostics
 - 5.3.3.1 Legrange Multiplier test
 - 5.3.3.2 Conditional SD plot
 - 5.3.3.3 Students-T QQ Plot

Modeling non-stationarity

```
In [1]: import numpy as np
import pandas as pd
import yfinance as yf
import matplotlib.pyplot as plt
import seaborn as sns
from arch.unitroot import ADF, KPSS
from arch import arch_model
import statsmodels.api as sm
from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
from statsmodels.tsa.api import ARIMA
from statsmodels.tsa.statespace.sarimax import SARIMAX
from statsmodels.stats.diagnostic import acorr_ljungbox
import scipy.stats as stats
from pmdarima import auto_arima
import warnings
import sys

warnings.filterwarnings("ignore")
sns.set(style='whitegrid', rc={"grid.linewidth": 0.1})
sns.set_context("paper")
params = {
    'legend.fontsize': 'medium',
    'axes.labelsize': 'large',
    'xtick.labelsize': 'medium',
    'ytick.labelsize': 'medium',
    'axes.titlesize': 'large',
    'figure.figsize': (12,9)
}
plt.rcParams.update(params)
```

Selecting the dataset

Microsoft, Tesla and Bitcoin's closing prices with daily frequencies were chosen from 2020 to 2021.

```
In [2]: dataset = yf.download(["MSFT", "TSLA", "BTC-USD"], start="2019-01-01", end="2022-01-01")["Adj Close"].dropna()  
dataset.head()
```


[*****100%*****] 3 of 3 completed

```
Out[2]:
```

	BTC-USD	MSFT	TSLA
Date			
2018-12-31	3742.700439	97.580017	22.186666
2019-01-02	3943.409424	97.147713	20.674667
2019-01-03	3836.741211	93.573822	20.024000
2019-01-04	3857.717529	97.925896	21.179333
2019-01-07	4025.248291	98.050781	22.330667

Plot chosen dataset

```
In [3]: plt.subplot(3, 1, 1)  
plt.plot(dataset['MSFT'])  
plt.title('Microsoft')  
  
plt.subplot(3, 1, 2)  
plt.plot(dataset['TSLA'])  
plt.title('Tesla')  
  
plt.subplot(3, 1, 3)  
plt.plot(dataset['BTC-USD'])  
plt.title('Bitcoin');
```



We can realize that all item have trend.

We can determine that ADF test will be executed with trend.

Stationary Test

ADF Test with trend

```
In [25]: microsoft_adf_with_trend = ADF(dataset['MSFT'], trend="ct")
tesla_adf_with_trend = ADF(dataset['TSLA'], trend="ct")
bitcoin_adf_with_trend = ADF(dataset['BTC-USD'], trend="ct")
print(f"Test statistics \n{microsoft_adf_with_trend.summary()}")
print(f"Test statistics \n{tesla_adf_with_trend.summary()}")
print(f"Test statistics \n{bitcoin_adf_with_trend.summary()}")
```

```
Test statistics
  Augmented Dickey-Fuller Results
=====
Test Statistic      -2.655
P-value              0.255
Lags                 16
-----
```

```
Trend: Constant and Linear Time Trend
Critical Values: -3.97 (1%), -3.42 (5%), -3.13 (10%)
Null Hypothesis: The process contains a unit root.
Alternative Hypothesis: The process is weakly stationary.
Test statistics
  Augmented Dickey-Fuller Results
=====
Test Statistic      -2.296
P-value              0.436
Lags                 0
-----
```

```
Trend: Constant and Linear Time Trend
Critical Values: -3.97 (1%), -3.42 (5%), -3.13 (10%)
Null Hypothesis: The process contains a unit root.
Alternative Hypothesis: The process is weakly stationary.
Test statistics
  Augmented Dickey-Fuller Results
=====
Test Statistic      -1.989
P-value              0.608
Lags                 12
-----
```

```
Trend: Constant and Linear Time Trend
Critical Values: -3.97 (1%), -3.42 (5%), -3.13 (10%)
Null Hypothesis: The process contains a unit root.
Alternative Hypothesis: The process is weakly stationary.
```

KPSS test

```
In [5]: microsoft_kpss = KPSS(dataset['MSFT'])
tesla_kpss = KPSS(dataset['TSLA'])
bitcoin_kpss = KPSS(dataset['BTC-USD'])
print(microsoft_kpss.summary())
print(tesla_kpss.summary())
print(bitcoin_kpss.summary())
```

```

      KPSS Stationarity Test Results
=====
Test Statistic          4.134
P-value                 0.000
Lags                    17
-----

Trend: Constant
Critical Values: 0.74 (1%), 0.46 (5%), 0.35 (10%)
Null Hypothesis: The process is weakly stationary.
Alternative Hypothesis: The process contains a unit root.
      KPSS Stationarity Test Results
=====
Test Statistic          3.902
P-value                 0.000
Lags                    17
-----

Trend: Constant
Critical Values: 0.74 (1%), 0.46 (5%), 0.35 (10%)
Null Hypothesis: The process is weakly stationary.
Alternative Hypothesis: The process contains a unit root.
      KPSS Stationarity Test Results
=====
Test Statistic          3.343
P-value                 0.000
Lags                    17
-----

Trend: Constant
Critical Values: 0.74 (1%), 0.46 (5%), 0.35 (10%)
Null Hypothesis: The process is weakly stationary.
Alternative Hypothesis: The process contains a unit root.
```

We have conclude that all items cannot reject the null hypothesis that the process contains a unit root, which means all item are not stationary by ADF with trend and KPSS test.

Mitigation of non-stationarity

```
In [6]: diff_data = dataset.diff().dropna()
diff_data.head()
```

Out[6]:

	BTC-USD	MSFT	TSLA
Date			
2019-01-02	200.708984	-0.432304	-1.511999
2019-01-03	-106.668213	-3.573891	-0.650667
2019-01-04	20.976318	4.352074	1.155333
2019-01-07	167.530762	0.124886	1.151335
2019-01-08	5.599609	0.710922	0.025999

ADF Test

```
In [7]: diff_microsoft_adf_with_trend = ADF(diff_data['MSFT'], trend="ct")
diff_tesla_adf_with_trend = ADF(diff_data['TSLA'], trend="ct")
diff_bitcoin_adf_with_trend = ADF(diff_data['BTC-USD'], trend="ct")
print(f"Test statistics \n{diff_microsoft_adf_with_trend.summary()}")
print(f"Test statistics \n{diff_tesla_adf_with_trend.summary()}")
print(f"Test statistics \n{diff_bitcoin_adf_with_trend.summary()}")
```

Test statistics

Augmented Dickey-Fuller Results

=====

Test Statistic	-8.128
----------------	--------

P-value	0.000
---------	-------

Lags	12
------	----

Trend: Constant and Linear Time Trend

Critical Values: -3.97 (1%), -3.42 (5%), -3.13 (10%)

Null Hypothesis: The process contains a unit root.

Alternative Hypothesis: The process is weakly stationary.

Test statistics

Augmented Dickey-Fuller Results

=====

Test Statistic	-27.683
----------------	---------

P-value	0.000
---------	-------

Lags	0
------	---

Trend: Constant and Linear Time Trend

Critical Values: -3.97 (1%), -3.42 (5%), -3.13 (10%)

Null Hypothesis: The process contains a unit root.

Alternative Hypothesis: The process is weakly stationary.

Test statistics

Augmented Dickey-Fuller Results

=====

Test Statistic	-8.025
----------------	--------

P-value	0.000
---------	-------

Lags	11
------	----

Trend: Constant and Linear Time Trend

Critical Values: -3.97 (1%), -3.42 (5%), -3.13 (10%)

Null Hypothesis: The process contains a unit root.

Alternative Hypothesis: The process is weakly stationary.

KPSS test

```
In [8]: diff_microsoft_kpss = KPSS(diff_data['MSFT'])
diff_tesla_kpss = KPSS(diff_data['TSLA'])
diff_bitcoin_kpss = KPSS(diff_data['BTC-USD'])
print(diff_microsoft_kpss.summary())
print(diff_tesla_kpss.summary())
print(diff_bitcoin_kpss.summary())
```

KPSS Stationarity Test Results

=====

Test Statistic	0.099
----------------	-------

P-value	0.589
---------	-------

Lags	8
------	---

Trend: Constant

Critical Values: 0.74 (1%), 0.46 (5%), 0.35 (10%)

Null Hypothesis: The process is weakly stationary.

Alternative Hypothesis: The process contains a unit root.

KPSS Stationarity Test Results

=====

Test Statistic	0.182
----------------	-------

P-value	0.306
---------	-------

Lags	4
------	---

Trend: Constant

Critical Values: 0.74 (1%), 0.46 (5%), 0.35 (10%)

Null Hypothesis: The process is weakly stationary.

Alternative Hypothesis: The process contains a unit root.

KPSS Stationarity Test Results

=====

Test Statistic	0.091
----------------	-------

P-value	0.633
---------	-------

Lags	6
------	---

Trend: Constant

Critical Values: 0.74 (1%), 0.46 (5%), 0.35 (10%)

Null Hypothesis: The process is weakly stationary.

Alternative Hypothesis: The process contains a unit root.

The null-hypothesis, that the process is stationary cannot be rejected.


The above tests suggests that a difference component should be added to the model.

Parameter selection for time series modeling

ACF and PACF plots of the chosen datasets (First difference)

```
In [9]: fig, ax = plt.subplots(3,2, figsize=(14,14))

plot_acf(diff_data["MSFT"], ax=ax[0][0], title="ACF plot of Microsoft")
plot_pacf(diff_data["MSFT"], ax=ax[0][1], method="ywm", title="PACF plot of microsoft");
plot_acf(diff_data["TSLA"], ax=ax[1][0], title="ACF plot of Tesla")
plot_pacf(diff_data["TSLA"], ax=ax[1][1], method="ywm", title="PACF plot of Tesla");
plot_acf(diff_data["BTC-USD"], ax=ax[2][0], title="ACF plot of Bitcoin")
plot_pacf(diff_data["BTC-USD"], ax=ax[2][1], method="ywm", title="PACF plot of Bitcoin");
```



We have concluded to choose Microsoft prices as the ACF and PACF plot clearly shows that Tesla and Bitcoin does not have any significant ARMA component that can be modeled.

Microsoft

Lag(1) is significant in both the ACF and PACF plots.

An AR(1) component and MA(1) component exists, which are both negative.

MA(2) is positive and significant.

There are some seasonal components on 7th and 9th lags

Modeling the series

From the above inferences, we can use an ARIMA(1,1,2) for Microsoft as initial parameters to model the time-series.

ARIMA(1,1,2)

```
In [10]: arima_112 = SARIMAX(dataset['MSFT'], order=(1,1,2)).fit()
```

```
In [11]: print(arima_112.summary())
```

```

                        SARIMAX Results
=====
Dep. Variable:          MSFT      No. Observations:          758
Model:                SARIMAX(1, 1, 2)  Log Likelihood        -2009.475
Date:                 Tue, 04 Oct 2022  AIC                   4026.950
Time:                 15:44:56      BIC                   4045.467
Sample:               0              HQIC                   4034.082
                        - 758
Covariance Type:      opg
=====
              coef      std err          z      P>|z|      [0.025      0.975]
-----
ar.L1         -0.0363      0.287      -0.127      0.899      -0.598      0.526
ma.L1         -0.1692      0.283      -0.597      0.550      -0.724      0.386
ma.L2          0.0896      0.064      1.402      0.161      -0.036      0.215
sigma2         11.8352      0.385     30.703      0.000     11.080     12.591
=====
Ljung-Box (L1) (Q):           0.09  Jarque-Bera (JB):           467.20
Prob(Q):                     0.76  Prob(JB):                 0.00
Heteroskedasticity (H):       5.61  Skew:                   -0.30
Prob(H) (two-sided):          0.00  Kurtosis:                6.80
=====
```

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

From the summary, no coefficient rejects the null hypothesis. So, the initial ARIMA parameters needs to be modified for a better fit.

Auto-ARIMA

auto_arima module from pmdarima selects the parameters that minimizes the selected information criterion (AIC in our case).

```
In [12]: auto_model = auto_arima(
    dataset["MSFT"],
    start_p=0,
    start_d=0,
    start_q=0,
    max_p=3,
    max_d=3,
    max_q=3,
    trace=True,
    with_intercept=False,
    return_valid_fits=True
)
```

```
Performing stepwise search to minimize aic
ARIMA(0,1,0)(0,0,0)[0]           : AIC=4058.795, Time=0.02 sec
ARIMA(1,1,0)(0,0,0)[0]           : AIC=4025.207, Time=0.02 sec
ARIMA(0,1,1)(0,0,0)[0]           : AIC=4029.518, Time=0.03 sec
ARIMA(2,1,0)(0,0,0)[0]           : AIC=4025.709, Time=0.04 sec
ARIMA(1,1,1)(0,0,0)[0]           : AIC=4026.097, Time=0.05 sec
ARIMA(2,1,1)(0,0,0)[0]           : AIC=4027.428, Time=0.11 sec
ARIMA(1,1,0)(0,0,0)[0] intercept : AIC=4017.841, Time=0.04 sec
ARIMA(0,1,0)(0,0,0)[0] intercept : AIC=4054.831, Time=0.02 sec
ARIMA(2,1,0)(0,0,0)[0] intercept : AIC=4019.060, Time=0.09 sec
ARIMA(1,1,1)(0,0,0)[0] intercept : AIC=4019.203, Time=0.10 sec
ARIMA(0,1,1)(0,0,0)[0] intercept : AIC=4021.803, Time=0.05 sec
ARIMA(2,1,1)(0,0,0)[0] intercept : AIC=4020.985, Time=0.19 sec
```

Best model: ARIMA(1,1,0)(0,0,0)[0] intercept

Total fit time: 0.750 seconds

Now that we have found ARIMA(1,1,0) minimizes the AIC criterion, we can use the parameters to fit an ARIMA model and diagnose the coefficients

```
In [13]: arima_110 = SARIMAX(dataset['MSFT'], order=(1,1,0)).fit()
```

```
In [14]: print(arima_110.summary())
```

```
SARIMAX Results
=====
Dep. Variable:          MSFT    No. Observations:          758
Model:                SARIMAX(1, 1, 0)    Log Likelihood          -2010.603
Date:                 Tue, 04 Oct 2022    AIC                   4025.207
Time:                 15:44:57            BIC                   4034.465
Sample:                0                HQIC                  4028.773
                        - 758
Covariance Type:        opg
=====
              coef    std err          z      P>|z|      [0.025     0.975]
-----
ar.L1         -0.2143     0.021    -10.268     0.000    -0.255    -0.173
sigma2         11.8709     0.377     31.504     0.000     11.132     12.609
=====
Ljung-Box (L1) (Q):           0.01    Jarque-Bera (JB):           505.17
Prob(Q):                     0.93    Prob(JB):                  0.00
Heteroskedasticity (H):       5.67    Skew:                      -0.36
Prob(H) (two-sided):          0.00    Kurtosis:                   6.94
=====

Warnings:
[1] Covariance matrix calculated using the outer product of gradients (complex-step).
```

True value vs fitted value

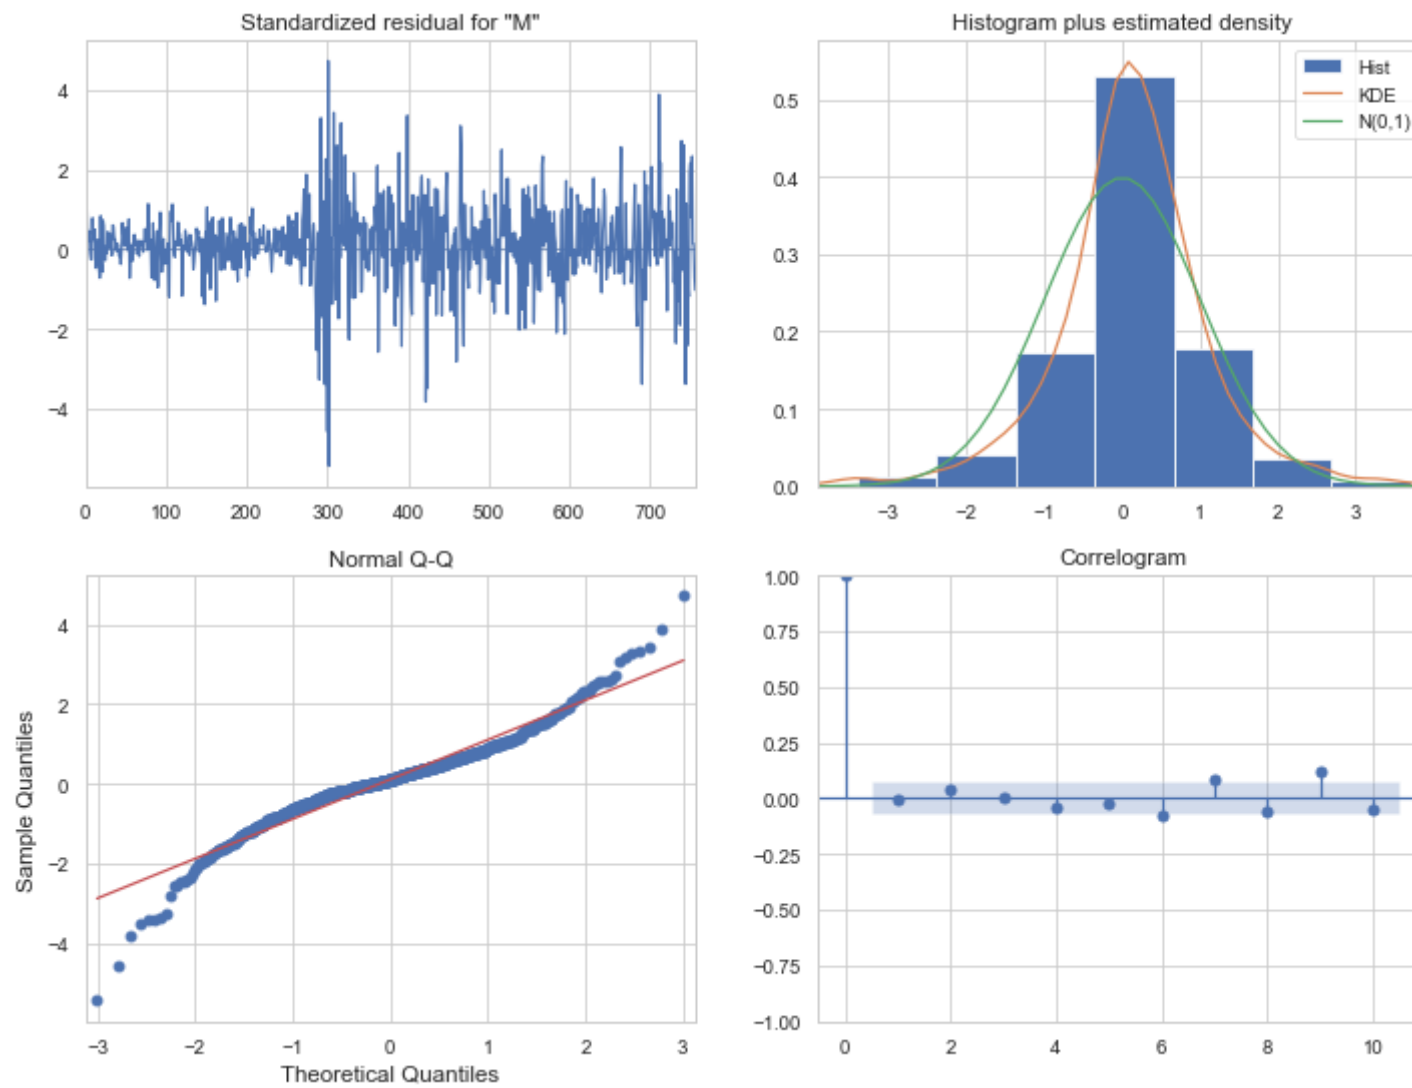
```
In [15]: plt.plot(dataset['MSFT'])
plt.plot(arima_110.fittedvalues[1:])
```

```
Out[15]: [<matplotlib.lines.Line2D at 0x21ef1043b20>]
```



Model diagnostics

```
In [16]: arima_110.plot_diagnostics();
```



Modeling heteroskedasticity in residuals with GARCH

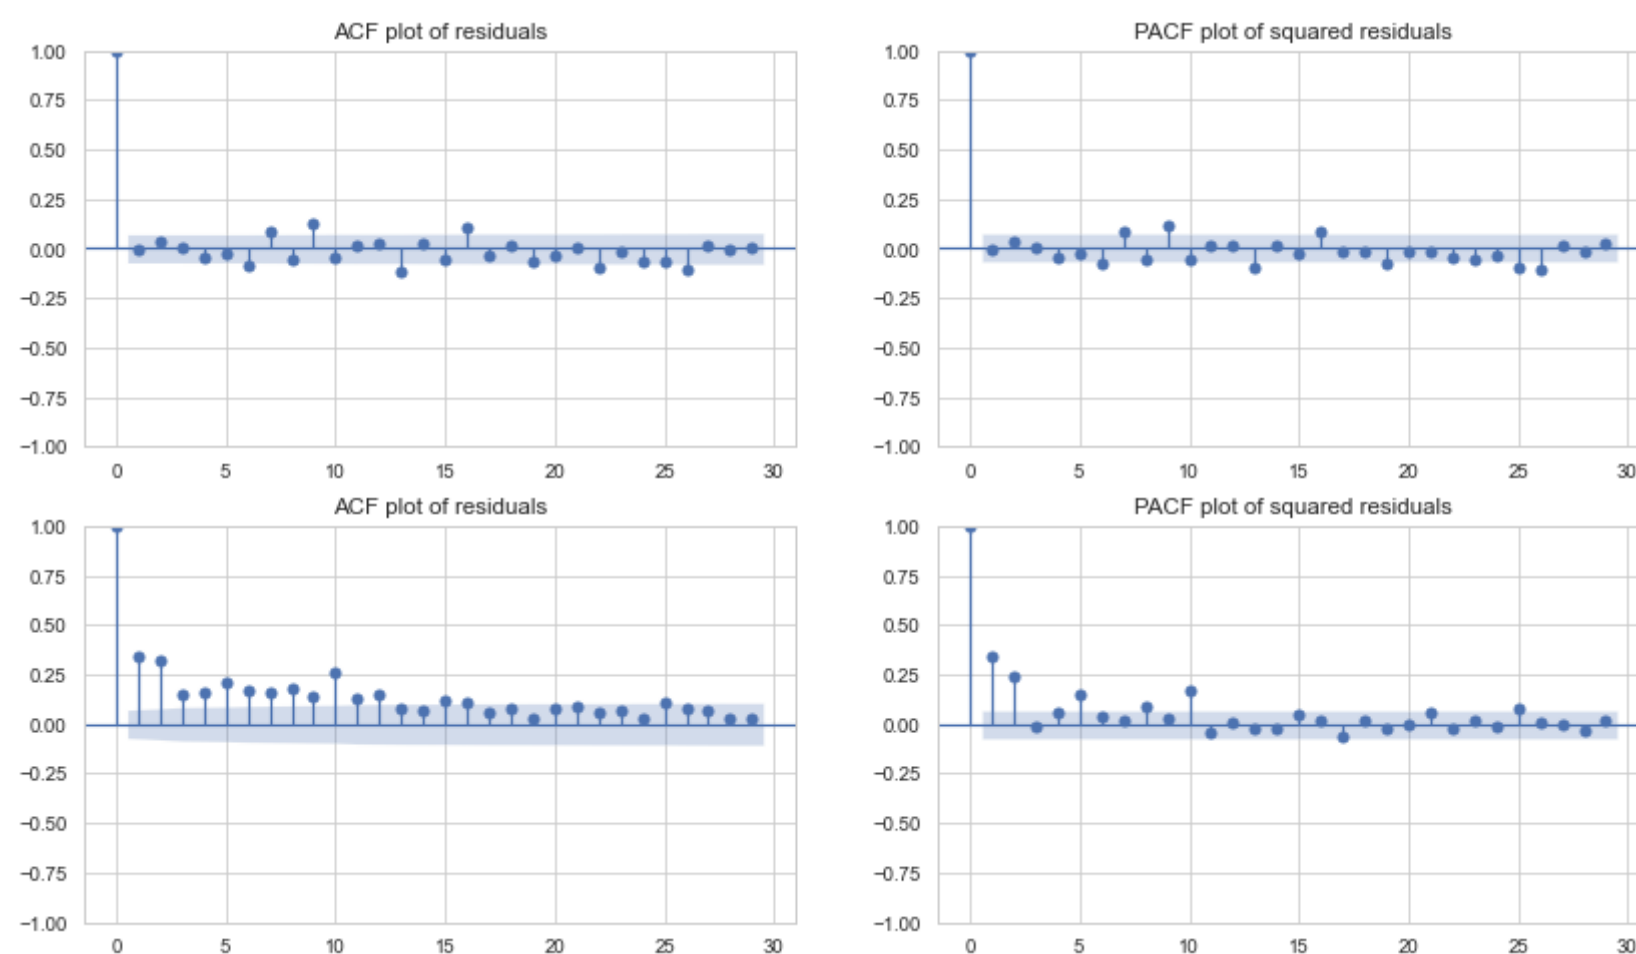
The residual plot shows clear heteroskedasticity, which suggests GARCH effects.

Also, the distribution is leptokurtic. So, a students-T distribution can be used for the GARCH model.

ACF and PACF of the residuals and the squares

```
In [17]: fig, ax = plt.subplots(2,2, figsize=(14,8))

plot_acf(arima_110.resid[1:], ax=ax[0][0], title="ACF plot of residuals")
plot_pacf(arima_110.resid[1:], ax=ax[0][1], method="yw", title="PACF plot of squared residuals");
plot_acf(arima_110.resid[1:] ** 2, ax=ax[1][0], title="ACF plot of residuals")
plot_pacf(arima_110.resid[1:] ** 2, ax=ax[1][1], method="yw", title="PACF plot of squared residuals");
```



GARCH fit

Let's begin with a GARCH(2,1) model.

```
In [18]: garch_21 = arch_model(arima_110.resid[1:], mean="constant", p=2, q=1, dist="studentsT").fit(update_freq=5)
```

```
Iteration:      5,   Func. Count:      41,   Neg. LLF: 1844.998766359014
Iteration:     10,   Func. Count:      80,   Neg. LLF: 1844.580363372783
Optimization terminated successfully   (Exit mode 0)
      Current function value: 1844.5802009055478
      Iterations: 12
      Function evaluations: 94
      Gradient evaluations: 12
```

```
In [19]: print(garch_21.summary())
```

```

              Constant Mean - GARCH Model Results
=====
Dep. Variable:              None    R-squared:                0.000
Mean Model:                Constant Mean    Adj. R-squared:          0.000
Vol Model:                 GARCH    Log-Likelihood:        -1844.58
Distribution:              Standardized Student's t    AIC:                3701.16
Method:                   Maximum Likelihood    BIC:                3728.94
                                     No. Observations:          757
Date:                     Tue, Oct 04 2022    Df Residuals:          756
Time:                     15:44:58    Df Model:              1
                                     Mean Model
=====
              coef      std err          t      P>|t|   95.0% Conf. Int.
-----
mu           0.4101   7.097e-02     5.778  7.574e-09 [ 0.271,  0.549]
              Volatility Model
=====
              coef      std err          t      P>|t|   95.0% Conf. Int.
-----
omega        0.0980   7.398e-02     1.325    0.185 [-4.696e-02,  0.243]
alpha[1]     0.1293   5.597e-02     2.310  2.091e-02 [1.957e-02,  0.239]
alpha[2]    7.9480e-03  6.578e-02     0.121    0.904 [ -0.121,  0.137]
beta[1]      0.8628   4.438e-02    19.442  3.384e-84 [ 0.776,  0.950]
              Distribution
=====
              coef      std err          t      P>|t|   95.0% Conf. Int.
-----
nu           6.6877     1.244     5.375  7.648e-08 [ 4.249,  9.126]
=====
```

Covariance estimator: robust

All the coefficients are statistically significant.

Diagnostics


Legrange Multiplier test

```
In [20]: # ARCH LM test for conditional heteroskedasticity
print("\nARCH LM test for conditional heteroskedasticity")
print(garch_21.arch_lm_test(standardized=True))
```

```
ARCH LM test for conditional heteroskedasticity
ARCH-LM Test
H0: Standardized residuals are homoskedastic.
ARCH-LM Test
H1: Standardized residuals are conditionally heteroskedastic.
Statistic: 9.6122
P-value: 0.9747
Distributed: chi2(20)
```

Conditional SD plot

```
In [21]: plt.plot(arima_110.resid[1:])
plt.plot(2*garch_21.conditional_volatility, c="r")
plt.plot(-2*garch_21.conditional_volatility, c="r")
plt.title("Residuals of ARIMA model with 2 Conditional SD");
```



Students-T QQ Plot

```
In [22]: sm.qqplot(garch_21.resid, stats.t, fit=True, line="q")
plt.title("Students-T QQ Plot");
```

